

Version Spaces + Candidate Elimination

Lecture Outline:

- Quick Review of Concept Learning and General-to-Specific Ordering
- Version Spaces
- The Candidate Elimination Algorithm
- Inductive Bias

Reading:

Chapter 2 of Mitchell

Version Spaces

- One limitation of the **FIND-S** algorithm is that it outputs just one hypothesis consistent with the training data – there might be many.

To overcome this, introduce notion of **version space** and algorithms to compute it.

Version Spaces

- One limitation of the **FIND-S** algorithm is that it outputs just one hypothesis consistent with the training data – there might be many.

To overcome this, introduce notion of **version space** and algorithms to compute it.

- A hypothesis h is **consistent** with a set of training examples D of target concept c if and only if $h(x) = c(x)$ for each training example $\langle x, c(x) \rangle$ in D .

$$\text{Consistent}(h, D) \equiv (\forall \langle x, c(x) \rangle \in D) h(x) = c(x)$$

Version Spaces

- One limitation of the **FIND-S** algorithm is that it outputs just one hypothesis consistent with the training data – there might be many.

To overcome this, introduce notion of **version space** and algorithms to compute it.

- A hypothesis h is **consistent** with a set of training examples D of target concept c if and only if $h(x) = c(x)$ for each training example $\langle x, c(x) \rangle$ in D .

$$\text{Consistent}(h, D) \equiv (\forall \langle x, c(x) \rangle \in D) h(x) = c(x)$$

- The **version space**, $VS_{H,D}$, with respect to hypothesis space H and training examples D , is the subset of hypotheses from H consistent with all training examples in D .

$$VS_{H,D} \equiv \{h \in H \mid \text{Consistent}(h, D)\}$$

Version Spaces

- One limitation of the **FIND-S** algorithm is that it outputs just one hypothesis consistent with the training data – there might be many.

To overcome this, introduce notion of **version space** and algorithms to compute it.

- A hypothesis h is **consistent** with a set of training examples D of target concept c if and only if $h(x) = c(x)$ for each training example $\langle x, c(x) \rangle$ in D .

$$\text{Consistent}(h, D) \equiv (\forall \langle x, c(x) \rangle \in D) h(x) = c(x)$$

- The **version space**, $VS_{H,D}$, with respect to hypothesis space H and training examples D , is the subset of hypotheses from H consistent with all training examples in D .

$$VS_{H,D} \equiv \{h \in H \mid \text{Consistent}(h, D)\}$$

- Note difference between definitions of **consistent** and **satisfies**:
 - an example x **satisfies** hypothesis h when $h(x) = 1$, regardless of whether x is +ve or –ve example of target concept
 - an example x is **consistent** with hypothesis h iff $h(x) = c(x)$

The LIST-THEN-ELIMINATE Algorithm

- Can represent version space by listing all members.
- Leads to **List-Then-Eliminate** concept learning algorithm:

1. $VersionSpace \leftarrow$ a list containing every hypothesis in H
2. For each training example, $\langle x, c(x) \rangle$
remove from $VersionSpace$ any hypothesis h for which $h(x) \neq c(x)$
3. Output the list of hypotheses in $VersionSpace$


- **List-Then-Eliminate** works in principle, so long as version space is finite.
- However, since it requires exhaustive enumeration of all hypotheses in practice it is not feasible.
- Is there a more compact way to represent version spaces?

The CANDIDATE-ELIMINATION Algorithm

- The **Candidate-Elimination** algorithm is similar to **List-Then-Eliminate** algorithm but uses a more compact representation of version space.
 - represents version space by its **most general** and **most specific** members

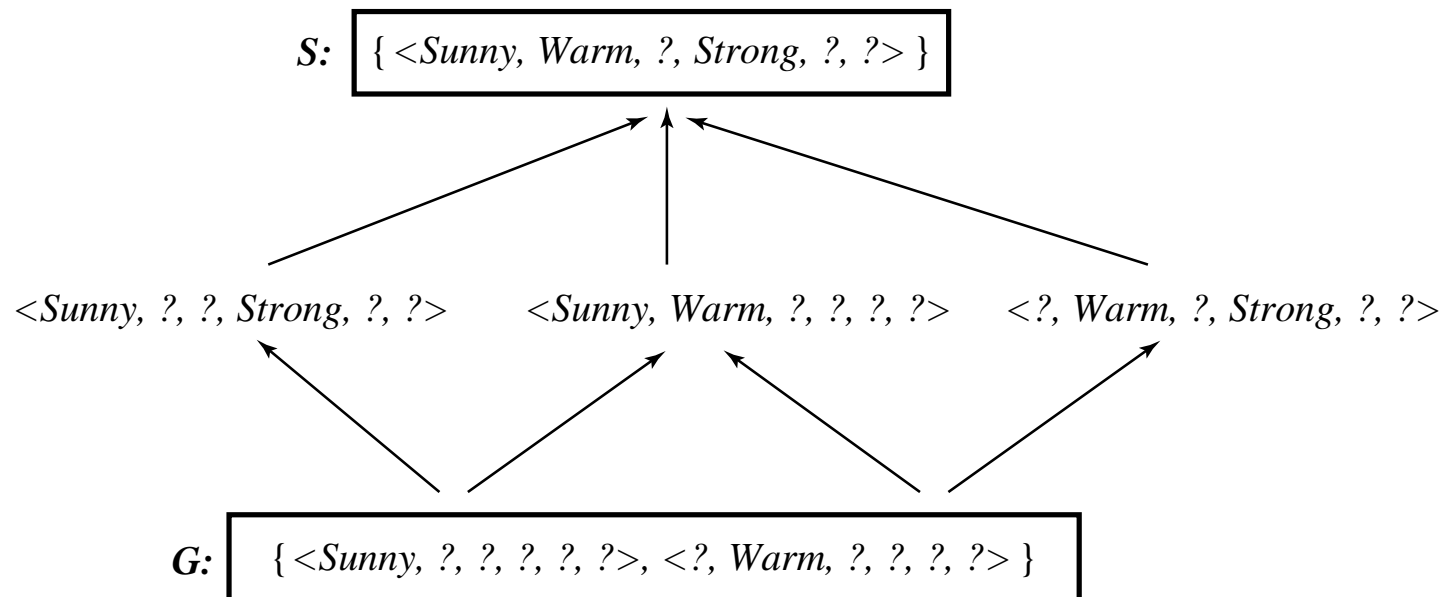
The CANDIDATE-ELIMINATION Algorithm

- The **Candidate-Elimination** algorithm is similar to **List-Then-Eliminate** algorithm but uses a more compact representation of version space.
 - represents version space by its **most general** and **most specific** members
- For *EnjoySport* example **Find-S** outputs the hypothesis: $h = \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, ?, ? \rangle$ which was one of 6 hypotheses consistent with the data.



The CANDIDATE-ELIMINATION Algorithm

- The **Candidate-Elimination** algorithm is similar to **List-Then-Eliminate** algorithm but uses a more compact representation of version space.
 - represents version space by its **most general** and **most specific** members
- For *EnjoySport* example **Find-S** outputs the hypothesis: $h = \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, ?, ? \rangle$ which was one of 6 hypotheses consistent with the data.



- The **Candidate-Elimination** algorithm represents the version space by recording only the most general members (G) and its most specific members (S)
 - other intermediate members in general-to-specific ordering can be generated as needed

The CANDIDATE-ELIMINATION Algorithm (cont)

- The **General boundary**, G , of version space $VS_{H,D}$ is the set of its maximally general members
- The **Specific boundary**, S , of version space $VS_{H,D}$ is the set of its maximally specific members

The CANDIDATE-ELIMINATION Algorithm (cont)

- The **General boundary**, G , of version space $VS_{H,D}$ is the set of its maximally general members
- The **Specific boundary**, S , of version space $VS_{H,D}$ is the set of its maximally specific members
- **Version Space Representation Theorem**

Every member of the version space lies between these boundaries

$$VS_{H,D} = \{h \in H \mid (\exists s \in S)(\exists g \in G)(g \geq_g h \geq_g s)\}$$

where $x \geq_g y$ means x is more general or equal to y

(see Mitchell, p. 32, for proof)

The CANDIDATE-ELIMINATION Algorithm (cont)

- The **General boundary**, G , of version space $VS_{H,D}$ is the set of its maximally general members
- The **Specific boundary**, S , of version space $VS_{H,D}$ is the set of its maximally specific members

- **Version Space Representation Theorem**

Every member of the version space lies between these boundaries

$$VS_{H,D} = \{h \in H \mid (\exists s \in S)(\exists g \in G)(g \geq_g h \geq_g s)\}$$

where $x \geq_g y$ means x is more general or equal to y

(see Mitchell, p. 32, for proof)

- Intuitively, **Candidate-Elimination** algorithm proceeds by
 - initialising G and S to the maximally general and maximally specific hypotheses in H
 - considering each training example in turn and
 - * using positive examples to drive the maximally specific boundary up
 - * using negative examples to drive the maximally general boundary down

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d
 - For each hypothesis s in S that is not consistent with d

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d
 - For each hypothesis s in S that is not consistent with d
 - * Remove s from S

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d
 - For each hypothesis s in S that is not consistent with d
 - * Remove s from S
 - * Add to S all minimal generalizations h of s such that

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d
 - For each hypothesis s in S that is not consistent with d
 - * Remove s from S
 - * Add to S all minimal generalizations h of s such that
 1. h is consistent with d , and
 2. some member of G is more general than h

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d
 - For each hypothesis s in S that is not consistent with d
 - * Remove s from S
 - * Add to S all minimal generalizations h of s such that
 1. h is consistent with d , and
 2. some member of G is more general than h
 - * Remove from S any hypothesis that is more general than another hypothesis in S

The CANDIDATE-ELIMINATION Algorithm (cont)

$G \leftarrow$ maximally general hypotheses in H

$S \leftarrow$ maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d
 - For each hypothesis s in S that is not consistent with d
 - * Remove s from S
 - * Add to S all minimal generalizations h of s such that
 1. h is consistent with d , and
 2. some member of G is more general than h
 - * Remove from S any hypothesis that is more general than another hypothesis in S
- If d is a negative example
 - Remove from S any hypothesis inconsistent with d
 - For each hypothesis g in G that is not consistent with d
 - * Remove g from G
 - * Add to G all minimal specializations h of g such that
 1. h is consistent with d , and
 2. some member of S is more specific than h
 - * Remove from G any hypothesis that is less general than another hypothesis in G

The CANDIDATE-ELIMINATION Algorithm: Example

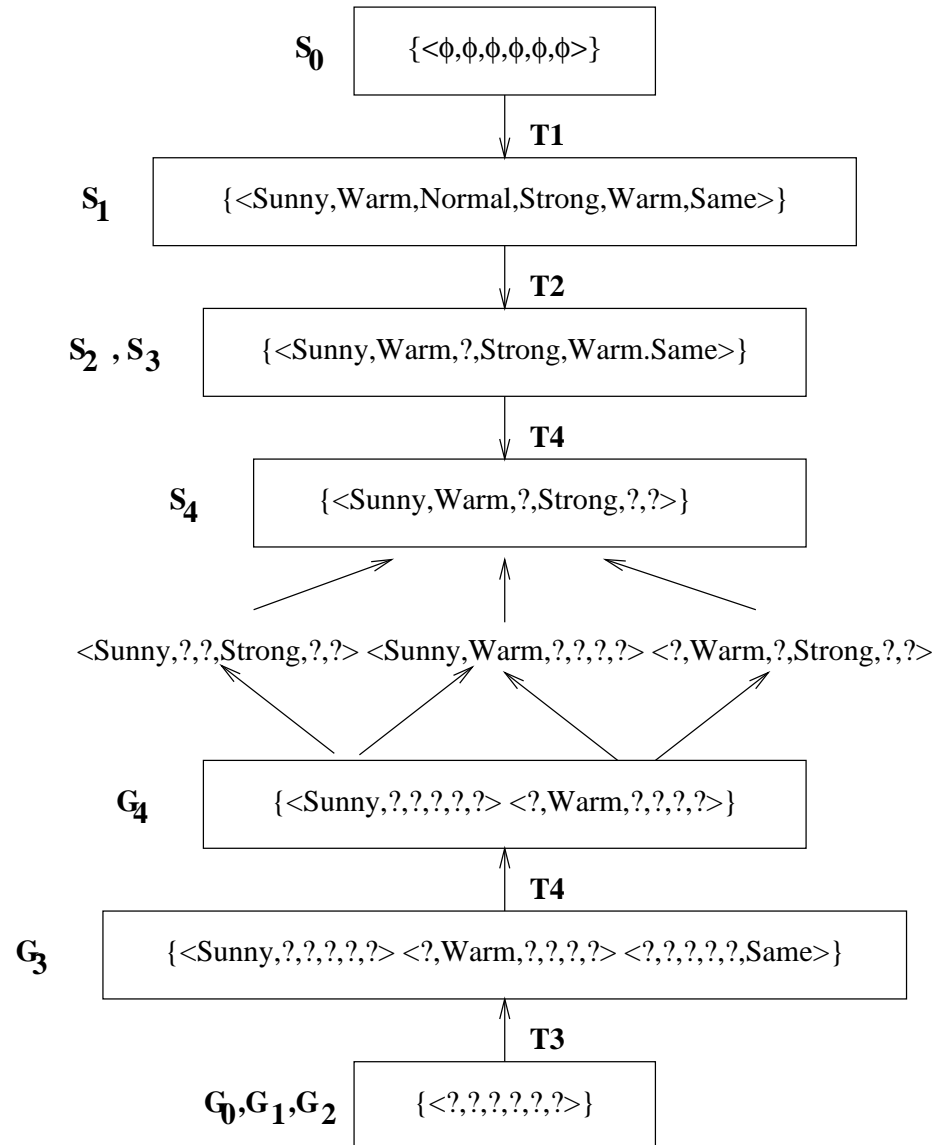
Training Examples:

T1: $\langle \text{Sunny}, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle, \text{Yes}$

T2: $\langle \text{Sunny}, \text{Warm}, \text{High}, \text{Strong}, \text{Warm}, \text{Same} \rangle, \text{Yes}$

T3: $\langle \text{Rainy}, \text{Cold}, \text{High}, \text{Strong}, \text{Warm}, \text{Change} \rangle, \text{No}$

T4: $\langle \text{Sunny}, \text{Warm}, \text{High}, \text{Strong}, \text{Cool}, \text{Change} \rangle, \text{Yes}$



The CANDIDATE-ELIMINATION Algorithm: Remarks

- Version space learned by **Candidate-Elimination** algorithm will converge towards correct hypothesis provided:
 - no errors in training examples
 - there is a hypothesis in H that describes target concept

In such cases algorithm may converge to empty version space

The CANDIDATE-ELIMINATION Algorithm: Remarks

- Version space learned by **Candidate-Elimination** algorithm will converge towards correct hypothesis provided:

- no errors in training examples
- there is a hypothesis in H that describes target concept

In such cases algorithm may converge to empty version space

- If algorithm can request next training example (e.g. from teacher) can increase speed of convergence by requesting examples that split the version space
 - E.g. T5: $\langle \text{Sunny, Warm, Normal, Light, Warm, Same} \rangle$ satisfies 3 hypotheses in previous example
 - * If T5 positive, S generalised, 3 hypotheses eliminated
 - * If T5 negative, G specialised, 3 hypotheses eliminated
 - Optimal query strategy is to request examples that exactly split version space – converge in $\lceil \log_2 |VS| \rceil$ steps. However, this is not always possible.

The CANDIDATE-ELIMINATION Algorithm: Remarks (cont)

- When using (i.e **not** training) a classifier that has not completely converged, new examples may be
 1. classed as positive by all $h \in VS$
 2. classed as negative by all $h \in VS$
 3. classed as positive by some, and negative by other, $h \in VS$

Cases 1 and 2 are unproblematic. In case 3. may want to consider proportion of positive vs. negative classifications (but then *a priori* probabilities of hypotheses are relevant)

Inductive Bias

- As noted, version space learned by **Candidate-Elimination** algorithm will converge towards correct hypothesis provided:
 - no errors in training examples
 - there is a hypothesis in H that describes target concept

What if no concept in H that describes the target concept?

Inductive Bias

- As noted, version space learned by **Candidate-Elimination** algorithm will converge towards correct hypothesis provided:
 - no errors in training examples
 - there is a hypothesis in H that describes target concept

What if no concept in H that describes the target concept?

- Consider the training data

Example	Sky	Temp	Humid	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Cloudy	Warm	Normal	Strong	Warm	Same	Yes
3	Rainy	Warm	Normal	Strong	Warm	Same	No

Inductive Bias

- As noted, version space learned by **Candidate-Elimination** algorithm will converge towards correct hypothesis provided:
 - no errors in training examples
 - there is a hypothesis in H that describes target concept

What if no concept in H that describes the target concept?

- Consider the training data

Example	Sky	Temp	Humid	Wind	Water	Forecast	EnjoySport
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Cloudy	Warm	Normal	Strong	Warm	Same	Yes
3	Rainy	Warm	Normal	Strong	Warm	Same	No

- No hypotheses consistent with 3 examples.

Most specific hypothesis consistent with Ex 1 and 2 *and representable in H* :

$\langle ?, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle$

But this is inconsistent with Ex 3.

Inductive Bias (cont)

- Need more expressive hypothesis representation language.

E.g. allow disjunctive or negative attribute values:

Sky = *Sunny* \vee *Cloudy*

Sky \neq *Rainy*

An Unbiased Learner

- What about ensuring **every** concept can be represented in H ?
 - Since concepts are subsets of instance space X , want H to be able to represent any set in power set of X
 - * for *EnjoySport* there were 96 possible instances
so, power set contains $2^{96} \approx 10^{28}$ possible target concepts
 - * recall biased conjunctive hypothesis space can represent only 973 of these
- Can do this by allowing hypotheses that are arbitrary conjunctions, disjunctions and negations of our earlier hypotheses
 - New problem: concept learning algorithm cannot generalise beyond observed examples!
 - * S boundary = disjunction of positive examples – exactly covers observed positive examples
 - * G boundary = negation of disjunction of negative examples – exactly rules out observed negative examples

An Unbiased Learner

- Capacity of **Candidate-Elimination** to generalise lies in its implicit assumption of **bias** – that target concept can be represented as a conjunction of attribute values

- Fundamental property of inductive inference:

a learner that makes no a priori assumptions regarding the identity of the target concept has no rational basis for classifying any unseen instances

I.e. **bias-free learning is futile**

Inductive Bias, More Formally

- Since all inductive learning involves bias, useful to characterise learning approaches by the type of bias they employ
- Consider
 - concept learning algorithm L
 - instances X , target concept c
 - training examples $D_c = \{\langle x, c(x) \rangle\}$
 - let $L(x_i, D_c)$ denote the classification, positive or negative, assigned to the instance x_i by L after training on data D_c .

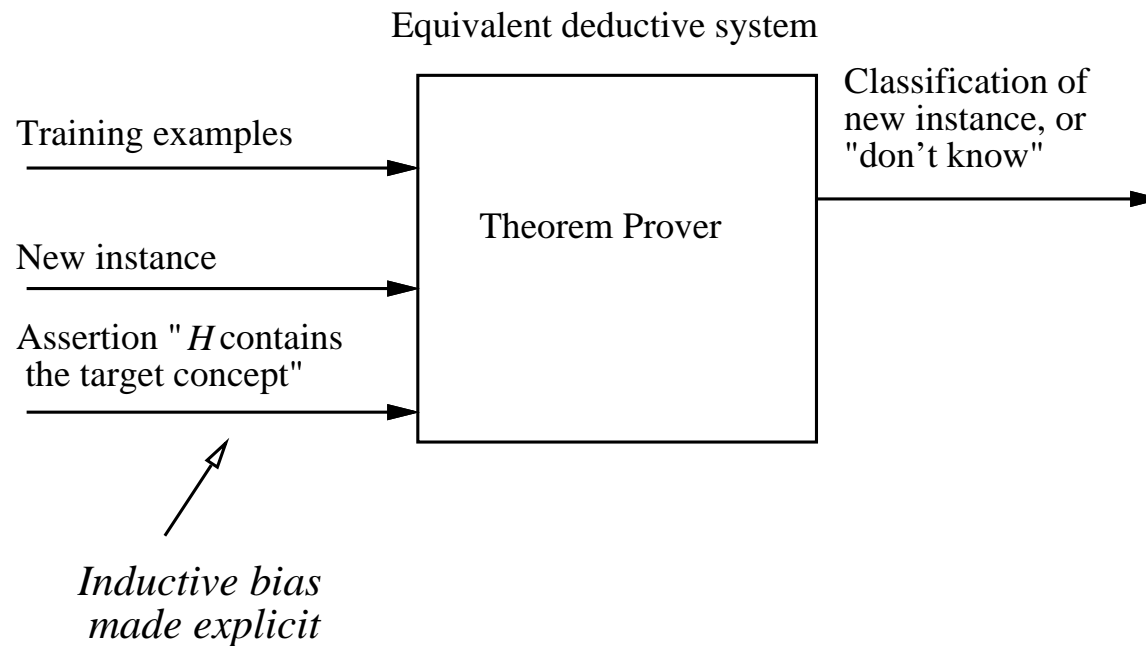
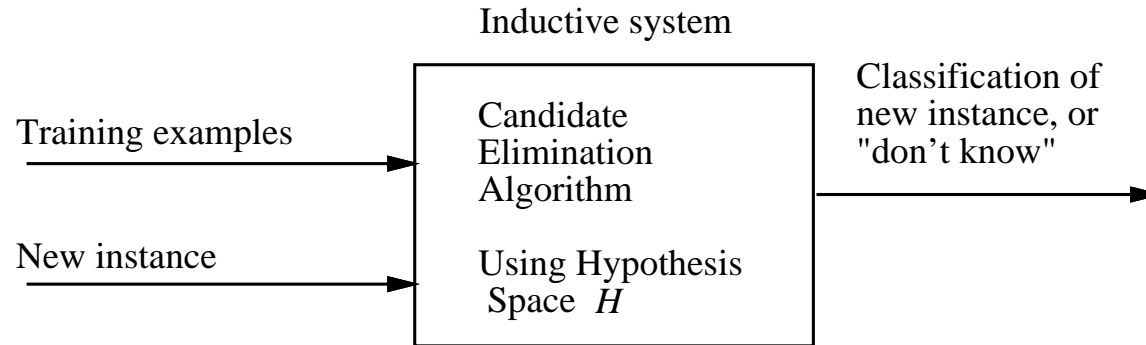
Definition:

The **inductive bias** of L is any minimal set of assertions B such that for any target concept c and corresponding training examples D_c

$$(\forall x_i \in X)[(B \wedge D_c \wedge x_i) \vdash L(x_i, D_c)]$$

where $A \vdash B$ means A logically entails B

Modelling Inductive Systems by Deductive Systems



Summary

- The **version space** with respect to a hypothesis space H and a set of training examples D is the subset of all hypotheses in H **consistent** with all the examples in D .

Summary

- The **version space** with respect to a hypothesis space H and a set of training examples D is the subset of all hypotheses in H **consistent** with all the examples in D .
- The version space may be compactly represented by recording its **general boundary** G and **specific boundary** S .

Every hypothesis in the version space is guaranteed to lie between G and S by the **version space representation theorem**.

Summary

- The **version space** with respect to a hypothesis space H and a set of training examples D is the subset of all hypotheses in H **consistent** with all the examples in D .
- The version space may be compactly represented by recording its **general boundary** G and **specific boundary** S .
Every hypothesis in the version space is guaranteed to lie between G and S by the **version space representation theorem**.
- The **Candidate-Elimination algorithm** exploits this theorem by searching for H for the version space by using the examples in training data D to progressively generalise the specific boundary and specialise the general boundary.

Summary

- The **version space** with respect to a hypothesis space H and a set of training examples D is the subset of all hypotheses in H **consistent** with all the examples in D .
- The version space may be compactly represented by recording its **general boundary** G and **specific boundary** S .
Every hypothesis in the version space is guaranteed to lie between G and S by the **version space representation theorem**.
- The **Candidate-Elimination algorithm** exploits this theorem by searching for H for the version space by using the examples in training data D to progressively generalise the specific boundary and specialise the general boundary.
- There are certain concepts the **Candidate-Elimination** algorithm cannot learn because of the **bias** of the hypothesis space – every concept must be representable as a conjunction of attribute values.

Summary

- The **version space** with respect to a hypothesis space H and a set of training examples D is the subset of all hypotheses in H **consistent** with all the examples in D .
- The version space may be compactly represented by recording its **general boundary** G and **specific boundary** S .
Every hypothesis in the version space is guaranteed to lie between G and S by the **version space representation theorem**.
- The **Candidate-Elimination algorithm** exploits this theorem by searching for H for the version space by using the examples in training data D to progressively generalise the specific boundary and specialise the general boundary.
- There are certain concepts the **Candidate-Elimination** algorithm cannot learn because of the **bias** of the hypothesis space – every concept must be representable as a conjunction of attribute values.
- In fact, all inductive learning supposes some *a priori* assumptions about the nature of the target concept, or else there is no basis for generalisation beyond observed examples: **bias-free learning is futile**.